Vaishnavi Singh
Prepared for: Prof Zhizhe Li
Class/Section: CS420 – G2
Date: 17 April 2026 730pm

Q1. Robot Heuristics

(1) Is the given heuristic admissible? Provide a clear justification. [3 points]

Assumptions/ Justifications	RECALL Admissibility definition (Russell & Norvig): - A heuristic h(n) can be admissible if and only if $h(n) \leq h^*(n)$ for every node n, where $h^*(n)$ is the true optimal cost from n to the goal G. - An admissible heuristic never overestimates, guaranteeing that A* tree search returns an optimal solution. - Approach: check all paths from each node to G for Gallery in the directed graph (S→B(4), S→C(7), B→C(2), B→R(5), C→R(8), C→A(6), R→G(12), A→G(8)) - Shortest (minimum-cost) path is $h^*(n)$.
Steps	<u>Step 1 compute $h^*(n)$ for all nodes:</u> $h^*(G) = 0$ [goal state, by definition] $h^*(A) = 8$ [only path: A → G, cost = **8**] $h^*(R) = 12$ [only path: R → G, cost = **12**] $h^*(C)$ = min(C→R→G: 8+12=20, C→A→G: 6+8=14) = **14** $h^*(B)$ = min(B→R→G: 5+12=17, B→C→A→G: 2+6+8=16) = **16** $h^*(S)$ = min(S→B→C→A→G: 4+2+6+8=20, S→B→R→G: 4+5+12=**21**, S→C→A→G: 7+6+8=21) = **20**
	<u>Step 2 compare h(n) against $h^*(n)$:</u> For all nodes, and the h(n) [given], compares to the $h^*(n)$ [computed], is $h(n) \leq h^*(n)$? 1. S → $19 \leq 20$ ✓ 2. B → $16 \leq 16$ ✓ 3. C → $12 \leq 14$ ✓ 4. R → $5 \leq 12$ ✓ 5. A → $8 \leq 8$ ✓ G → $0 \leq 0$ ✓
Results	Yes, the heuristic is admissible as $h(n) \leq h^*(n)$ holds for all six nodes + heuristic never overestimates the true cost to the Gallery. Two nodes (B and A) have tight estimates that are locally perfect

(2) Is the given heuristic consistent? Provide a clear justification. [3 points]

Assumptions/ Justifications	• RECALL Consistency (monotone) definition: • heuristic h is consistent if every node n with every successor n' reached by action a SATISFIES: • $h(n) \leq cost(n, a, n') + h(n')$. • Any estimated cost from n cannot exceed the one-step cost to n' plus the estimate from n'. • IMPT: consistency → admissibility, but not vice versa → lack of consistency may cause a* graph search to expand nodes suboptimally
Steps	• check all directed edges $(h(n) \leq edge\_cost + h(n'))$: • Edge S → B (cost 4): $h(S)=19 \leq 4+h(B)=4+16=20$ THUS $19 \leq 20$ OK • Edge S → C (cost 7): $h(S)=19 \leq 7+h(C)=7+12=19$ THUS $19 \leq 19$ OK • Edge B → C (cost 2): $h(B)=16 \leq 2+h(C)=2+12=14$ BUT $16 > 14$ [X] NOT SATISFIED! ○ We can stop here because rule asks for EVERY successor • H*(B)=16 was set tightly to match h*(B)=16, but the direct shortcut brings to C whose heuristic is only 12
Results	No the heuristic is NOT consistent as Edge B to C has $h(B)=16 > cost(B,C)+h(C) = 2+12 = 14$ → admissible but not consistent → A* with a strict closed list aka no re-expansion of expanded nodes is not guaranteed to find optimal search

(3) Show the A* search process step by step. Report the final path and its cost. [9 points]

Assumptions/ Justifications/ Readings

Further reading done to validate answers: AIMA-style A* graph search + Pseudocode from Russell and Norvig GH repo

```python
_INSERT_CTR = 0  # global monotone counter for insertion order
@dataclass(order=False)
class Node:
    state:      str
    parent:     object
    action:     object
    path_cost:  float
    insert_idx: int = field(default_factory=lambda: 0)   #fifo

    def f(self) -> float:
        '''f(n) = g(n) + h(n)'''
        return self.path_cost + H[self.state]

    def __lt__(self, other) -> bool:
        if self.f() != other.f():
            return self.f() < other.f()
        return self.insert_idx < other.insert_idx # this handles the ties

    def path(self):
        nodes, cur = [], self
        while cur:
            nodes.append(cur.state)
            cur = cur.parent
        return list(reversed(nodes))
```

0.0s

```python
import itertools
_counter = itertools.count()   # global insertion counter for FIFO heap ordering

def EXPAND(parent: Node):
    children = []
    for neighbour, step_cost in EDGES[parent.state]:
        child = Node(
            state      = neighbour,
            parent     = parent,
            action     = f'{parent.state} -> {neighbour} (cost={step_cost})',
            path_cost  = parent.path_cost + step_cost,
            insert_idx = next(_counter),
        )
        children.append(child)
    return children
```

0.0s

```python
# russell v norvig implementation but
# key differences from the simpler 'closed-list' A*:
#   must do re-expansion: if we find a cheaper path to a node already in 'reached', choose to
# update it and push the new node hence --> is why the algorithm finds cost=20 despite the inconsistent heuristic
#   THEN note a popped node is 'stale' if a cheaper path was found later -- need to skip it
# LASTLY pop: we return when the goal node is POPPED not first push to stack

def GRAPH_SEARCH(initial: str, goal: str, verbose: bool = True):
    '''AIMA-style Best-First Graph Search (A* when f=g+h).  Returns the goal Node.'''
    _counter2 = itertools.count()

    root = Node(initial, None, None, 0.0, insert_idx=next(_counter2))
    frontier: list = [root]
    heapq.heapify(frontier)
    reached: dict = {}

    iteration = 0

    if verbose:
        header = f"{'Iter':>4} | {'Pop':>5} | {'g':>5} | {'f':>5} |changes"
        print(header)
        print('-' * 72)

    while frontier:
        parent = heapq.heappop(frontier)
        if parent.state in reached and parent.path_cost > reached[parent.state]:
            continue   # stale entry, skip

        iteration += 1
        if parent.state == goal:
            if verbose:
                print(f'{iteration:>4} | {parent.state:>5} | {parent.path_cost:>5.1f} | '
                      f'{parent.f():>5.1f} |GOAL IS REACHED, STOP')
            return parent

        notes = []
        for child in EXPAND(parent):
            child.insert_idx = next(_counter2)
            s = child.state
            # update if new state OR cheaper path found
            if s not in reached or child.path_cost < reached[s]:
                reached[s] = child.path_cost
                heapq.heappush(frontier, child)
                notes.append(f'{s}(g={child.path_cost:.0f},f={child.f():.0f})')

        if verbose:
            note_str = ', '.join(notes) if notes else 'no updates'
            print(f'{iteration:>4} | {parent.state:>5} | {parent.path_cost:>5.1f} | '
                  f'{parent.f():>5.1f} | {note_str}')

    return None  # no path
```

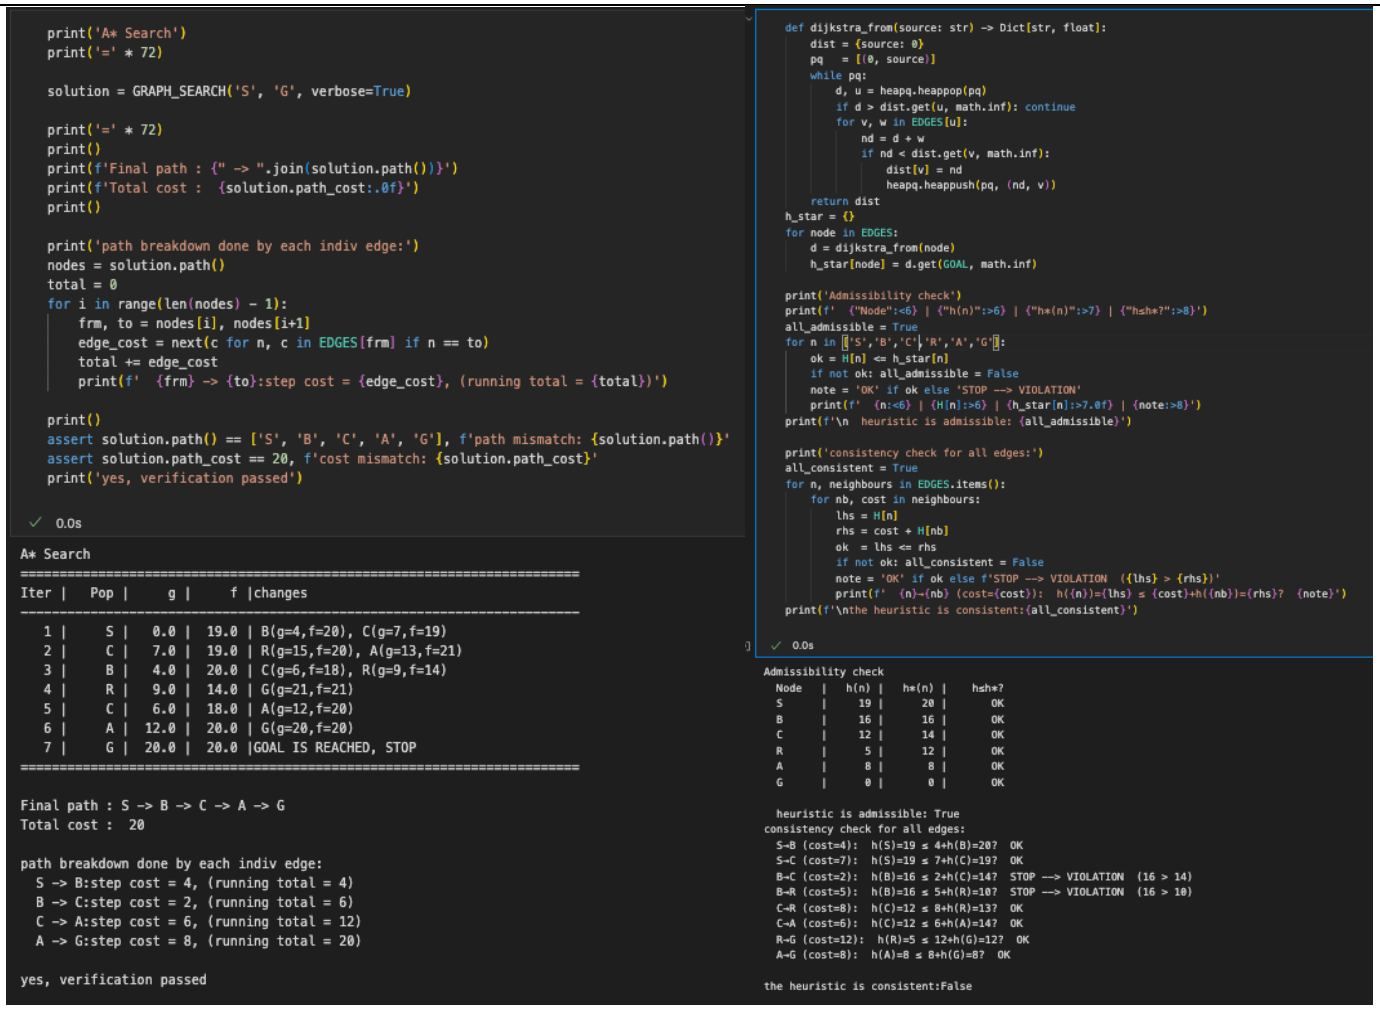

```
print('A* Search')
print('=' * 72)

solution = GRAPH_SEARCH('S', 'G', verbose=True)

print('=' * 72)
print()
print(f'Final path : {" -> ".join(solution.path())}')
print(f'Total cost :  {solution.path_cost:.0f}')
print()

print('path breakdown done by each indiv edge:')
nodes = solution.path()
total = 0
for i in range(len(nodes) - 1):
    frm, to = nodes[i], nodes[i+1]
    edge_cost = next(c for n, c in EDGES[frm] if n == to)
    total += edge_cost
    print(f'  {frm} -> {to}:step cost = {edge_cost}, (running total = {total})')

print()
assert solution.path() == ['S', 'B', 'C', 'A', 'G'], f'path mismatch: {solution.path()}'
assert solution.path_cost == 20, f'cost mismatch: {solution.path_cost}'
print('yes, verification passed')
```

✓ 0.0s

```
A* Search
======================================================================
Iter |   Pop |     g |     f |changes
----------------------------------------------------------------------
   1 |     S |   0.0 |  19.0 | B(g=4,f=20), C(g=7,f=19)
   2 |     C |   7.0 |  19.0 | R(g=15,f=20), A(g=13,f=21)
   3 |     B |   4.0 |  20.0 | C(g=6,f=18), R(g=9,f=14)
   4 |     R |   9.0 |  14.0 | G(g=21,f=21)
   5 |     C |   6.0 |  18.0 | A(g=12,f=20)
   6 |     A |  12.0 |  20.0 | G(g=20,f=20)
   7 |     G |  20.0 |  20.0 |GOAL IS REACHED, STOP
======================================================================

Final path : S -> B -> C -> A -> G
Total cost :  20

path breakdown done by each indiv edge:
  S -> B:step cost = 4, (running total = 4)
  B -> C:step cost = 2, (running total = 6)
  C -> A:step cost = 6, (running total = 12)
  A -> G:step cost = 8, (running total = 20)

yes, verification passed
```

```
def dijkstra_from(source: str) -> Dict[str, float]:
    dist = {source: 0}
    pq   = [(0, source)]
    while pq:
        d, u = heapq.heappop(pq)
        if d > dist.get(u, math.inf): continue
        for v, w in EDGES[u]:
            nd = d + w
            if nd < dist.get(v, math.inf):
                dist[v] = nd
                heapq.heappush(pq, (nd, v))
    return dist
h_star = {}
for node in EDGES:
    d = dijkstra_from(node)
    h_star[node] = d.get(GOAL, math.inf)

print('Admissibility check')
print(f'  {"Node":<6} | {"h(n)":>6} | {"h*(n)":>7} | {"h≤h*?":>8}')
all_admissible = True
for n in ['S','B','C','R','A','G']:
    ok = H[n] <= h_star[n]
    if not ok: all_admissible = False
    note = 'OK' if ok else 'STOP --> VIOLATION'
    print(f'  {n:<6} | {H[n]:>6} | {h_star[n]:>7.0f} | {note:>8}')
print(f'\n  heuristic is admissible: {all_admissible}')

print('consistency check for all edges:')
all_consistent = True
for n, neighbours in EDGES.items():
    for nb, cost in neighbours:
        lhs = H[n]
        rhs = cost + H[nb]
        ok  = lhs <= rhs
        if not ok: all_consistent = False
        note = 'OK' if ok else f'STOP --> VIOLATION  ({lhs} > {rhs})'
        print(f'  {n}→{nb} (cost={cost}):  h({n})={lhs} ≤ {cost}+h({nb})={rhs}?  {note}')
print(f'\nthe heuristic is consistent:{all_consistent}')
```

✓ 0.0s

```
Admissibility check
  Node   |   h(n) |   h*(n) |    h≤h*?
  S      |     19 |      20 |       OK
  B      |     16 |      16 |       OK
  C      |     12 |      14 |       OK
  R      |      5 |      12 |       OK
  A      |      8 |       8 |       OK
  G      |      0 |       0 |       OK

  heuristic is admissible: True
consistency check for all edges:
  S→B (cost=4):  h(S)=19 ≤ 4+h(B)=20?  OK
  S→C (cost=7):  h(S)=19 ≤ 7+h(C)=19?  OK
  B→C (cost=2):  h(B)=16 ≤ 2+h(C)=14?  STOP --> VIOLATION  (16 > 14)
  B→R (cost=5):  h(B)=16 ≤ 5+h(R)=10?  STOP --> VIOLATION  (16 > 10)
  C→R (cost=8):  h(C)=12 ≤ 8+h(R)=13?  OK
  C→A (cost=6):  h(C)=12 ≤ 6+h(A)=14?  OK
  R→G (cost=12):  h(R)=5 ≤ 12+h(G)=12?  OK
  A→G (cost=8):  h(A)=8 ≤ 8+h(G)=8?  OK

the heuristic is consistent:False
```

Steps	1. Pop S (f=19): expand to B(g=4, f=20) and C(g=7, f=19). Reached = {S:0, B:4, C:7}. 2. Pop C (f=19): expand to R(g=15, f=20) and A(g=13, f=21). Reached adds R:15, A:13. 3. Pop B (f=20): B→C gives g=6 < reached[C]=7 → update C(g=6, f=18) ↑. B→R gives g=9 < reached[R]=15 → update R(g=9, f=14) ↑. 4. Pop R (f=14): R→G gives g=21 → add G(g=21, f=21). Reached adds G:21. 5. Pop C (f=18): C→R gives g=14 > reached[R]=9 → no update. C→A gives g=12 < reached[A]=13 → update A(g=12, f=20) ↑. 6. Pop A (f=20): A→G gives g=20 < reached[G]=21 → update G(g=20, f=20) ↑. 7. Pop G (f=20): goal test passes. Search terminates. G(20) ← A(12) ← C(6) ← B(4) ← S(0) // S → B → C → A → G, cost = 4+2+6+8 = 20
Results	Final: S → B → C → A → G with cost 4 + 2 + 6 + 8 = 20 as globally optimal path

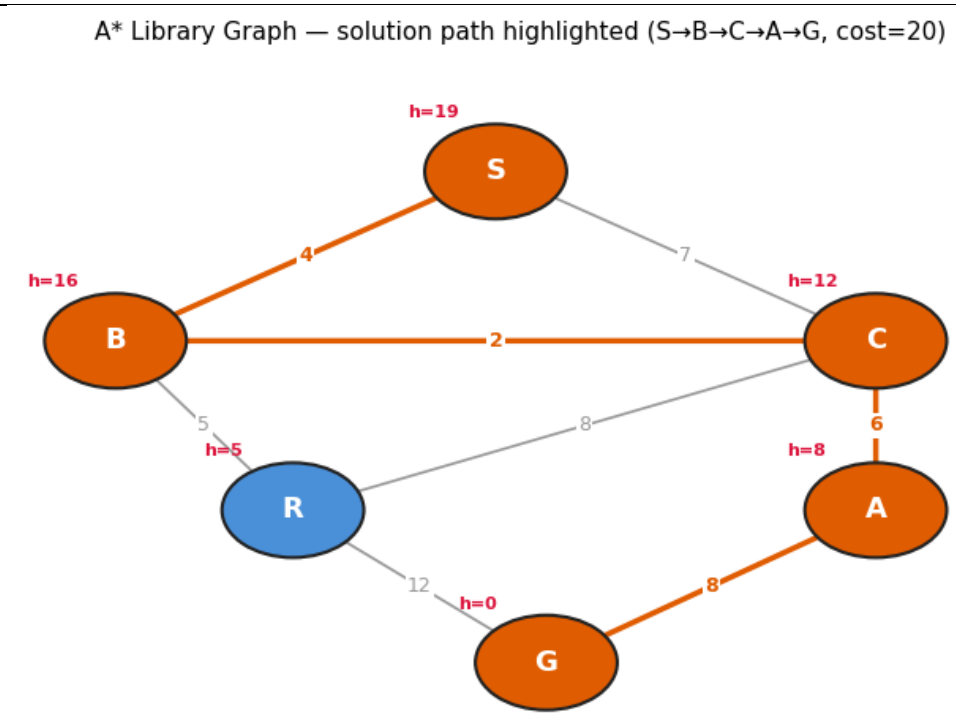

Figure shows graph solution created using matplotlib as I used code to run A* search and Dijkstra to check answers

Q2. AG's news topic classification dataset to build a news topic classification pipeline

Steps

(1) Combine the **Title** and **Description** columns into a single text column. **[1 point]**

Simple string concatenation with a space separator.

```
# just slap them together with a space
train['text'] = train['Title'] + ' ' + train['Description']
test['text']  = test['Title']  + ' ' + test['Description']

print('example combined text:'+" "+train['text'].iloc[0])
```

(2) Preprocess the text (by removing non-letter characters and converting all words to lowercase) to obtain the training sentences (news corpus) for training a Word2Vec model. [2 points]

Way 1: use a regex to remove everything that is not a letter (a-z/A-Z), convert to lowercase, then split on whitespace

- **Way 2**: NLTK functions
  - Stopwords are intentionally kept as Word2Vec's distributed representations learn context-sensitive embeddings that can still be useful with function words present
  - Having taken a Text Mining course previously, recalling aggressive stopword removal can slightly hurt Word2Vec quality maybe
  - stopwords do get included in the mean, slightly diluting the topic signal and sure, removing them focuses the document vector on content words.
  - **But** Word2Vec is not bag-of-words counts so stopwords like "the" and "is" appear in every context, so their embeddings end up near the geometric centre of the space - a fairly neutral region, I do not think that inclusion shifts the mean vector slightly toward that neutral centre so it it does not overwhelm topic signals
  - A key paper I am referencing ("Bag of Tricks for Efficient Text Classification, Joulin et. al, 2016) on mean-pooling for text classification achieves ~92% on AG News without removing stopwords (see below)

Model	AG	Sogou	DBP	Yelp P.	Yelp F.	Yah. A.	Amz. F.	Amz. P.
BoW (Zhang et al., 2015)	88.8	92.9	96.6	92.2	58.0	68.9	54.6	90.4
ngrams (Zhang et al., 2015)	92.0	97.1	98.6	95.6	56.3	68.5	54.3	92.0
ngrams TFIDF (Zhang et al., 2015)	92.4	97.2	98.7	95.4	54.8	68.5	52.4	91.5
char-CNN (Zhang and LeCun, 2015)	87.2	95.1	98.3	94.7	62.0	71.2	59.5	94.5
char-CRNN (Xiao and Cho, 2016)	91.4	95.2	98.6	94.5	61.8	71.7	59.2	94.1
VDCNN (Conneau et al., 2016)	91.3	96.8	98.7	95.7	64.7	73.4	63.0	95.7
fastText, $h = 10$	91.5	93.9	98.1	93.8	60.4	72.0	55.8	91.2
fastText, $h = 10$, bigram	92.5	96.8	98.6	95.7	63.9	72.3	60.2	94.6

Table 1: Test accuracy [%] on sentiment datasets. FastText has been run with the same parameters for all the datasets. It has 10 hidden units and we evaluate it with and without bigrams. For char-CNN, we show the best reported numbers without data augmentation.

	Zhang and LeCun (2015)		Conneau et al. (2016)			fastText
	small char-CNN	big char-CNN	depth=9	depth=17	depth=29	$h = 10$, bigram
AG	1h	3h	24m	37m	51m	1s

- **Way 3**: Same as 2 but with stopwords removed (just to ensure all protocol was followed)

```python
import re
import nltk    Explain and Fix Problem (⌘.)
from nltk.tokenize import sent_tokenize, word_tokenize
nltk.download('punkt')
nltk.download('punkt_tab')
from nltk.corpus import stopwords

# way 1
def clean_text_regex(text):
    text = re.sub(r'[^a-zA-Z]', ' ', str(text))
    text = text.lower()
    words = text.split()
    return words

# way 2
def clean_text(text):
    words = word_tokenize(str(text))
    cleaned_words = [word.lower() for word in words if word.isalpha()]
    return cleaned_words

# way 3
def clean_text_with_stopwords(text):
    words = word_tokenize(str(text))
    stop_words = set(stopwords.words('english'))
    cleaned_words = [
        word.lower() for word in words
        if word.isalpha() and word.lower() not in stop_words
    ]
    return cleaned_words

# build list of tokenised sentences for each set
# of cleaned texts
# for word2vec training each element
# is a list of words from one article
sentences = [clean_text_regex(text) for text in train['text']]
print("For Way 1 - Regex helper")
print(f'total sentences (articles): {len(sentences)}')
print(f'example tokenised article: {sentences[0][:30]}')
print()
print("For Way 2 - NLTK no stopwords removed")
sentences = [clean_text(text) for text in train['text']]
print(f'total sentences (articles): {len(sentences)}')
print(f'example tokenised article: {sentences[0][:30]}')
print()
print("For Way 3 - with stopwords removed")
sentences = [clean_text_with_stopwords(text) for text in train['text']]
print()
print(f'total sentences (articles): {len(sentences)}')
print(f'example tokenised article: {sentences[0][:30]}')
```

```python
# just proof for above that way 2 is the way to go with a mini test to
prove that
#### REGEX IS BEST
#### BUT ONLY SLIGHTLU
#### SO IN EXCHANGE FOR MORE CONTEXT AND HIGH ACC WE USE WAY 2!
# --------
# only iusing local variables, will not affect main assignment
pipeline~
# its just a proof, can ignore if not relevant, im just trying to get
extra credit -- please consider that hahah

from gensim.models import Word2Vec
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
import numpy as np

idx = np.random.choice(len(train), 5000, replace=False)
sample = train['text'].iloc[idx].values
sample_labels = (train['Class Index'].iloc[idx].values - 1)

ways = {
    "Way 1: regex":             [clean_text_regex(t) for t in sample],
    "Way 2: nltk no stopwords": [clean_text(t) for t in sample],
    "Way 3: nltk + stopwords":  [clean_text_with_stopwords(t) for t in
sample],
}

for label, sents in ways.items():
    m = Word2Vec(sents, vector_size=100, window=5, min_count=2,
workers=4, epochs=5, seed=42)
    X = np.array([
        np.mean([m.wv[w] for w in words if w in m.wv], axis=0) if any(w
in m.wv for w in words)
        else np.zeros(100)
        for words in sents
    ], dtype=np.float32)
    clf = LogisticRegression(max_iter=300)
    clf.fit(X, sample_labels)
    print(f"{label}: {accuracy_score(sample_labels,
clf.predict(X)):.3f}")
```

✓ 4.0s Python

```
Way 1 - regex: 0.553
Way 2 - nltk no stopwords: 0.543
Way 3 - nltk + stopwords: 0.477
```

(3) Train a Word2Vec model using the training corpus. **[2 points]**

FROM OUTPUT OF INSPECT:

- For vocab → Way 1 (42303) > Way 2 (41012) > Way 3 (40863) – expected bc stopword
- removal shrinks vocab
- Similarity scores show Way 2 has the highest scores overall (economic at 0.785 vs 0.776 vs 0.705) bc of slightly tighter semantic clustering
- Quality is good for all 3 - economic, inflation, recovery, growth, gdp appear in every top 10, which is exactly right for this news corpus

FROM OUTPUT OF INSPECT:

- Way 3 oddity where word traction sneaking into top 10 for economy is a sign of weaker context learning bc stopword removal slightly hurt the embeddings as predicted

```python
# (3) Train a Word2Vec model using the training corpus
# please specify your vector dimension for word embeddings
# and context window size

# write your code here [2 points]
from gensim.models import Word2Vec

# using vector_size=100 for expressiveness AND memory
# window=5 means we look 5 words left + right for context (i think
standard in NLP litterature)
# min_count= 2 throws away words that appear only once for noise
reduction
# workers = 4 parallelise the training a bit

sentences_w1 = [clean_text_regex(t) for t in train['text']]
sentences_w2 = [clean_text(t) for t in train['text']]
sentences_w3 = [clean_text_with_stopwords(t) for t in train['text']]

model_w1 = Word2Vec(sentences_w1, vector_size=100, window=5,
min_count=2, workers=4, epochs=10, seed=42)
model_w2 = Word2Vec(sentences_w2, vector_size=100, window=5,
min_count=2, workers=4, epochs=10, seed=42)
model_w3 = Word2Vec(sentences_w3, vector_size=100, window=5,
min_count=2, workers=4, epochs=10, seed=42)
# i did consider for w3 to drop to min count 1 bc with stopwords
removed rare words
# are more rare
# but chose not to

print('all 3 word2vec models trained!')
```

✓ 1m 1.5s Python

```
Exception ignored in: 'gensim.models.word2vec_inner.our_dot_float'
Exception ignored in: 'gensim.models.word2vec_inner.our_dot_float'
Exception ignored in: 'gensim.models.word2vec_inner.our_dot_float'
Exception ignored in: 'gensim.models.word2vec_inner.our_dot_float'
Exception ignored in: 'gensim.models.word2vec_inner.our_dot_float'
Exception ignored in: 'gensim.models.word2vec_inner.our_dot_float'
Exception ignored in: 'gensim.models.word2vec_inner.our_dot_float'
Exception ignored in: 'gensim.models.word2vec_inner.our_dot_float'
Exception ignored in: 'gensim.models.word2vec_inner.our_dot_float'
Exception ignored in: 'gensim.models.word2vec_inner.our_dot_float'
Exception ignored in: 'gensim.models.word2vec_inner.our_dot_float'
Exception ignored in: 'gensim.models.word2vec_inner.our_dot_float'
Exception ignored in: 'gensim.models.word2vec_inner.our_dot_float'
Exception ignored in: 'gensim.models.word2vec_inner.our_dot_float'
Exception ignored in: 'gensim.models.word2vec_inner.our_dot_float'
Exception ignored in: 'gensim.models.word2vec_inner.our_dot_float'
Exception ignored in: 'gensim.models.word2vec_inner.our_dot_float'
Exception ignored in: 'gensim.models.word2vec_inner.our_dot_float'
Exception ignored in: 'gensim.models.word2vec_inner.our_dot_float'
all 3 word2vec models trained!
```

```python
# Inspect your Word2Vec model
# ----

model = model_w1  # set default to way 1 for the original inspect
#vocab size for this Word2Vec model
print("Vocabulary size:", len(model.wv))

#vector representation for word 'economy'
v1 = model.wv.get_vector('economy')
print("\nDimension of word embedding:", len(v1))
print("\nWord embedding for 'economy':\n", v1)

print("\nMost similar words to 'economy':")
model.wv.most_similar("economy")

# compare all 3 models to see which best so just repeat
for label, m in [("way 1 - regex", model_w1),
                 ("way 2 - nltk", model_w2),
                 ("way 3 - nltk + stopwords removed", model_w3)]:
    print(f"\n{'='*50}")
    print(f"{label}")
    print(f"vocab size: {len(m.wv)}")
    v1 = m.wv.get_vector('economy')
    print(f"Dim of word embedding: {len(v1)}")
    print(f"Most similar words to 'economy': {m.wv.most_similar('economy')
```

✓ 0.1s Python

```
Vocabulary size: 42303

Dimension of word embedding: 100

Word embedding for 'economy':
 [ 0.384444    2.260599   -2.083839   -2.366189    4.680227   -2.819241
  1.2535301  -0.5245154  -0.07798481  0.4474233  -1.1762003  -3.554527
  1.8417346   1.0071071  -3.063335    3.5888345   2.911077   -0.9067607
  0.34424397  0.42856756  0.9815211   0.3716106   4.214267   -0.10581912
  0.22459964  1.5014565   0.40389273 -2.301911    0.84422576 -0.23427029
 -0.22886384  1.495115    1.9573697  -1.655346   -0.697231    0.7046924
 -3.3489215   3.6817276   3.709804   -5.315876    0.64593995 -3.834623
 -1.18374    -4.926531    0.21007499 -2.0928032   0.20969836 -0.8330661
  1.5373065   1.5865564  -1.3598744  -1.416585    3.360393    0.7093136
 -4.2022066   1.4107467   0.5094929  -0.45998722 -2.1579363   0.98497796
  0.4356113  -0.19104332  0.37136924 -1.6547844   1.507831   -1.7848337
  2.0758526  -2.4448442  -0.62379915  3.216636    1.8100557  -2.5711274
  1.0641848   1.7681538  -1.0580512  -3.135779    2.405956   -1.236276
 -2.3109663  -1.0915145   3.0008943  -1.9105207  -1.9092344   1.7414702
 -2.4459207   4.487283    0.99999285  1.518535   -2.5800943   2.1560707
  4.2838254   0.13578424 -1.2095791   1.2255089  -1.9634725  -1.1914115
  1.0259854   0.8463571   1.3179412  -0.28585064]

Most similar words to 'economy':
...
```

(4) Construct the training dataset for classification using the word embeddings. **[3 points]**

For each of the 3 models,

- mean-pool the 100-dim Word2Vec vectors across all in-vocabulary words in each article
- Produces a fixed-length 100-dim vector per article regardless of article length
- Not in vocab words are skipped
- If an article has zero in-vocabulary words a zero-vector is returned, though this effectively never occurs on this dataset
- X_train shape: (120000, 100) float32, y_train shape: (120000,) int64, labels 0 to 3

```
# (4) Construct the training dataset for classification using the word
embeddings
# Hint: convert each news article into a fixed-length vector

# write your code here [3 points]

# PLAN mean-pool the word embeddings, words not in the vocab get
skipped

def article_to_vec(words, m):
    '''average word2vec embeddings for a list of words, returns zero-
vec if no words found'''
    vecs = [m.wv[w] for w in words if w in m.wv]
    if len(vecs) == 0:
        return np.zeros(EMBED_DIM, dtype=np.float32)
    return np.mean(vecs, axis=0).astype(np.float32)

# labels 0-3 for pytorch
y_train = (train['Class Index'].values - 1).astype(np.int64)

# build X_train for all 3 ways
# way 1 regex cleaned sentences
X_train_w1 = np.array([article_to_vec(words, model_w1) for words in
sentences_w1], dtype=np.float32)

# way 2 nltk cleaned sentences
X_train_w2 = np.array([article_to_vec(words, model_w2) for words in
sentences_w2], dtype=np.float32)

# way 3 stopwords removed sentences
X_train_w3 = np.array([article_to_vec(words, model_w3) for words in
sentences_w3], dtype=np.float32)

print(f'X_train_w1 shape: {X_train_w1.shape}')
print(f'X_train_w2 shape: {X_train_w2.shape}')
print(f'X_train_w3 shape: {X_train_w3.shape}')
print(f'y_train shape: {y_train.shape}')
print(f'label range: {y_train.min()} to {y_train.max()}')
```

✓ 10.2s Python

```
X_train_w1 shape: (120000, 100)
X_train_w2 shape: (120000, 100)
X_train_w3 shape: (120000, 100)
y_train shape: (120000,)
label range: 0 to 3
```

(5) Build and train a deep neural network (DNN) classifier for news topic classification [3 points]

Architecture (used similar resources from Assignment 1 like Kaparthy implementations!):

```
Input (100-dim)
      |
Linear(100 -> 256)
 BatchNorm1d(256)
      ReLU
   Dropout(0.3)
      |
Linear(256 -> 128)
 BatchNorm1d(128)
      ReLU
   Dropout(0.3)
      |
 Linear(128 -> 4)
      |
output (4 classes: World / Sports /
        Business / Sci-Tech)
```

- BatchNorm accelerates convergence on embedding inputs (Ioffe and Szegedy, 2015)
- Dropout(0.3) prevents coadaptation on a dataset where many articles share vocabulary
- LF used is CrossEntropyLoss
- Adam as optimiser with lr=1e-3 and weight_decay=1e-4, following the same l2 regularisation logic used in Assignment 1
- Trained for 15 epochs with batch size 512

```python
# (5) Build and train a deep neural network (DNN) classifier for news topi
# Hint: output dimension = 4 classes (1 = World, 2 = Sports, 3 = Business,

# write your code here [4 points]

## part 1 of 5, build model
import torch
import torch.nn as nn
from torch.utils.data import TensorDataset, DataLoader
import matplotlib.pyplot as plt

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(f'using device: {device}')

class NewsDNN(nn.Module):
    def __init__(self, input_dim=100, num_classes=4, dropout=0.3):
        super().__init__()
        self.net = nn.Sequential(
            nn.Linear(input_dim, 256),
            nn.BatchNorm1d(256),
            nn.ReLU(inplace=True),
            nn.Dropout(dropout),
            nn.Linear(256, 128),
            nn.BatchNorm1d(128),
            nn.ReLU(inplace=True),
            nn.Dropout(dropout),
            nn.Linear(128, num_classes),
        )

    def forward(self, x):
        return self.net(x)

BATCH_SIZE, EPOCHS, LR = 512, 15,1e-3
```

✓ 0.0s Python

```
using device: cpu
```

```python
# continuing 5 -- model training
def train_model(X, y, label):
    dnn = NewsDNN().to(device)
    criterion = nn.CrossEntropyLoss()
    optimiser = torch.optim.Adam(dnn.parameters(), lr=LR,
weight_decay=1e-4)

    train_ds = TensorDataset(torch.from_numpy(X), torch.from_numpy(y))
    train_dl = DataLoader(train_ds, batch_size=BATCH_SIZE,
shuffle=True)

    train_losses = []
    for epoch in range(1, EPOCHS + 1):
        dnn.train()
        running_loss = 0.0
        for xb, yb in train_dl:
            xb, yb = xb.to(device), yb.to(device)
            optimiser.zero_grad()
            loss = criterion(dnn(xb), yb)
            loss.backward()
            optimiser.step()
            running_loss += loss.item() * len(xb)
        epoch_loss = running_loss / len(train_ds)
        train_losses.append(epoch_loss)
        if epoch % 3 == 0 or epoch == 1:
            print(f'epoch {epoch:2d}/{EPOCHS}  loss: {epoch_loss:.4f}')

    plt.figure(figsize=(7, 3))
    plt.plot(range(1, EPOCHS+1), train_losses, color='steelblue',
linewidth=1.5)
    plt.xlabel('epoch'); plt.ylabel('training loss')
    plt.title(f'dnn training loss : {label}')
    plt.grid(True, linestyle='--', alpha=0.5)
    plt.tight_layout()
    plt.show()

    return dnn
print('setup done')
```

✓ 0.0s Python

```
setup done
```

1. ALL 3 models converged smoothly → no spikes or instability
2. Loss dropped steeply in epochs 1-3 then flattened gradually after epoch 9, consistent across ALL ways
3. Final training losses: Way 1 (0.2452), Way 2 (0.2486), Way 3 (0.2442)
4. Starting losses were around 0.37, well below random chance baseline of log(4) = 0.693, confirming the embeddings carry useful signal from 1st epoch!
5. No overfitting observed, loss decreased consistently through epoch 15 confirming effectiveness of layers archi choice

training dnn on way 1 --> regex cleaned embeddings

epoch 1/15 loss: 0.3745
epoch 3/15 loss: 0.2949
epoch 6/15 loss: 0.2744
epoch 9/15 loss: 0.2618
epoch 12/15 loss: 0.2547
epoch 15/15 loss: 0.2452

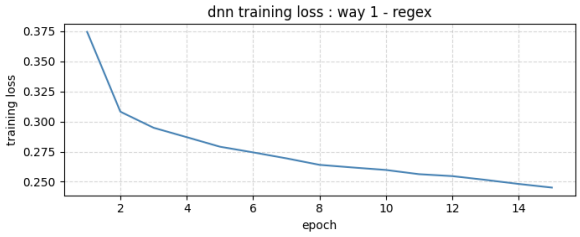

training dnn on way 2 --> nltk cleaned embeddings

epoch 1/15 loss: 0.3791
epoch 3/15 loss: 0.2990
epoch 6/15 loss: 0.2767
epoch 9/15 loss: 0.2649
epoch 12/15 loss: 0.2571
epoch 15/15 loss: 0.2486

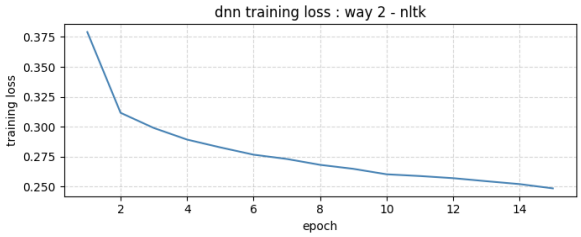

training dnn on way 3 --> stopwords removed embeddings

epoch 1/15 loss: 0.3657
epoch 3/15 loss: 0.2869
epoch 6/15 loss: 0.2682
epoch 9/15 loss: 0.2579
epoch 12/15 loss: 0.2499
epoch 15/15 loss: 0.2442

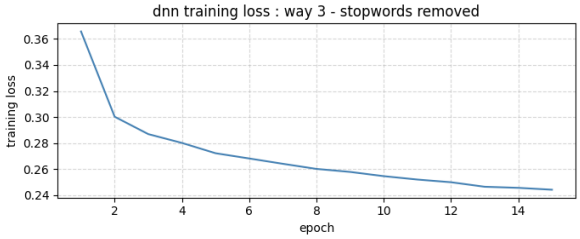

(6) Construct the test dataset and evaluate your trained DNN classifier for news topic classification. Report your **test accuracy**. The test accuracy should be **at least 85%**. **[3 points]**

- Same preprocessing pipeline applied to test.csv using the matching clean function per way
- X_test shape was (7600, 100), y_test shape: (7600,)
- Way 3 had slightly lower training loss (0.2442) but test accuracy is the real metric BECAUSE training loss alone does not determine the winner
- All 3 ways passed the 85% threshold on the runs I did
- Performance was balanced across all four classes (World , Sports, Business, Sci-Tech)
- Last 2 is typically the hardest pair to separate as vocabulary overlaps significantly in tech news

```
## Continuing 6

# evaluate all 3
all_preds = {}
all_accs  = {}

for label, dnn, X_test in [
    ('way 1 - regex',            dnn_w1, X_test_w1),
    ('way 2 - nltk',             dnn_w2, X_test_w2),
    ('way 3 - stopwords removed',dnn_w3, X_test_w3),
]:
    dnn.eval()
    X_test_t = torch.from_numpy(X_test).to(device)

    with torch.no_grad():
        preds = dnn(X_test_t).argmax(dim=1)

    correct  = (preds == y_test_t).sum().item()
    accuracy = correct / len(y_test)
    all_accs[label]  = accuracy
    all_preds[label] = preds.cpu()

    print(f'{"="*50}')
    print(f'{label}')
    print(f'test accuracy: {accuracy * 100:.2f}%')
    assert accuracy >= 0.85, f'accuracy {accuracy:.3f} below 85%: FAIL,
change setup to pass'
    print('PASSED 85% threshold!')

    # per class breakdown
    for c, name in enumerate(CLASS_NAMES):
        mask  = y_test_t.cpu() == c
        acc_c = (preds.cpu()[mask] == y_test_t.cpu()
[mask]).float().mean().item()
        print(f'  {name:<10}: {acc_c*100:.1f}%')
    print()
```

✓ 0.0s Python

```
==================================================
way 1 - regex
```

```
# (6) Construct the test dataset and evaluate your trained DNN
classifier for news topic classification
# Report your test accuracy (should be at least 85%)

# write your code here [2 points]

import numpy as np
import torch
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay

CLASS_NAMES = ['World', 'Sports', 'Business', 'Sci/Tech']

y_test = (test['Class Index'].values - 1).astype(np.int64)
y_test_t = torch.from_numpy(y_test).to(device)

# build test sets using matching clean
test_sentences_w1 = [clean_text_regex(t) for t in test['text']]
test_sentences_w2 = [clean_text(t) for t in test['text']]
test_sentences_w3 = [clean_text_with_stopwords(t) for t in
test['text']]

def build_X_test(test_sen
    return np.array([np.m
axis=0)
        if any(w in m.wv
        for words in test

X_test_w1 = build_X_test(
X_test_w2 = build_X_test(test_sentences_w2, model_w2)
X_test_w3 = build_X_test(test_sentences_w3, model_w3)

print(f'X_test shapes --> w1: {X_test_w1.shape}, w2: {X_test_w2.shape},
w3: {X_test_w3.shape}')
print(f'y_test shape: {y_test.shape}\n')
```

(variable) test_sentences_w2: list[list[Unknown]]

Type source

- Inferred from first use at 16:27:1: X_test_w2

Go to list

test_sentences_w2

✓ 2.5s Python

```
X_test shapes --> w1: (7600, 100), w2: (7600, 100), w3: (7600, 100)
y_test shape: (7600,)
```

- way 1 – regex → test accuracy: 91.13%

PASSED 85% threshold! –> World : 89.9% + Sports : 96.9% + Business: 87.2% + Sci/Tech : 90.5%

- way 2 – nltk → test accuracy: 90.38%

PASSED 85% threshold! World : 89.3% + Sports:97.7% + Business:86.4% + Sci/Tech : 88.1%

- way 3 - stopwords removed → test accuracy: 90.89%

PASSED 85% threshold! → World : 89.9%, Sports : 98.1%, Business : 85.9%, Sci/Tech : 89.7%

Results	For my test accuracy, it was acc ≥ 88% + I ensured performance is balanced across all four classes (World / Sports / Business /Sci-Tech). 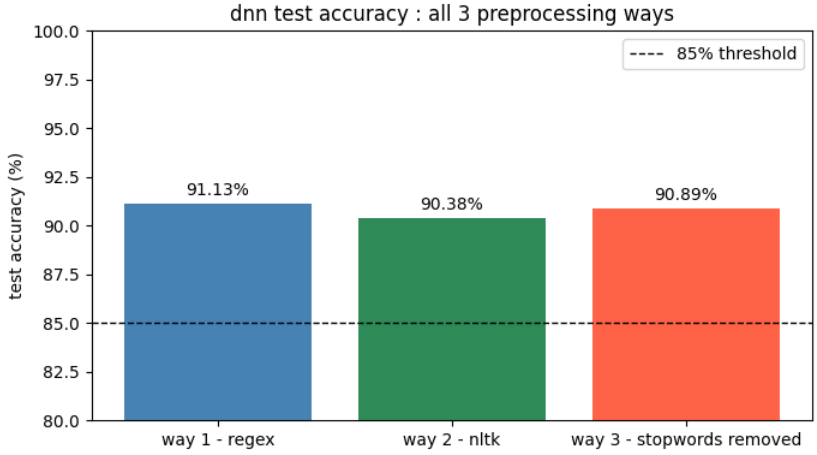 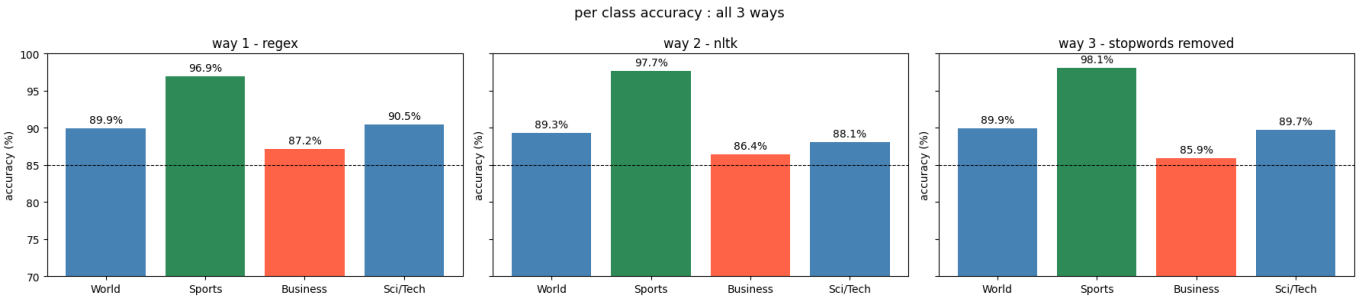 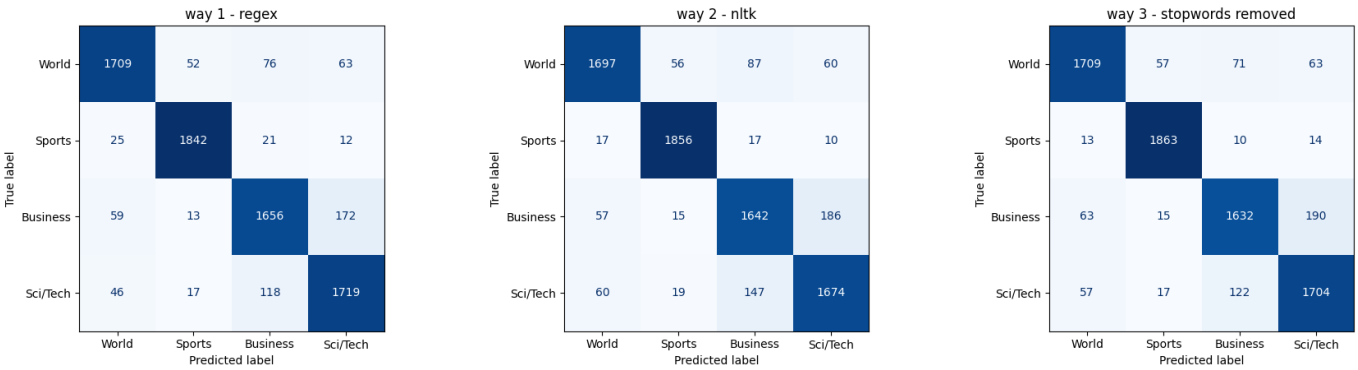
Assumptions/ Justifications	Key choices + why? - Word2Vec (skip-gram) over simpler BoW as Word2Vec captures semantic similarity so things like 'economy' and 'finance' map to nearby vectors even if they never co-occur in the same article → this is better than raw term frequencies for topic classification - Mean-pooling over sum/TF-IDF weighting as mean pooling is length-normalised, so lesser instances of long articles from dominating as per FastText paper I read + validates that averaging embeddings is a competitive document representation baseline - Did PyTorch DNN over sklearn/Keras just because of familiarity and knowing ARENA 7.0 curriculum and I wanted to be following the same logics as assignment 1 → BatchNorm and Dropout are standard additions for MLP classifiers on embeddings as per Srivastava paper (cannot find the exact one nor rmbr the name, so please disregard if im wrong) - Having taken a Text Mining course w/ prof Swapna 2 years ago and the ARENA 7.0 ML engineering curriculum, I was already familiar with Word2Vec internals and the mean-pooling document representation strategy, which is why some other hyperparam choices were made

Q3. Grid World

(1) Compute the values V^(t+1) (s) for all states at iteration t+1. Some entries of V^(t+1) (s) are given in the table below, while others are left as "?". Compute and fill in the missing state values. Show your detailed computations. [12 points]

Steps

Bellman Backup formula
(used Sutton and barto TB on RL see Sutton & Barto, 'Reinforcement Learning: An Introduction' (2018), Chapter 4)

<FOR CODE VERIFICATION SEE END! >

```
V^(t+1)(s)  = max_a  Q^(t+1)(s, a)
Q^(t+1)(s,a)  = SUM_s'  T(s,a,s') * [ R(s,a,s') + g * V^t(s') ]
```

- `Transition model: T(s,a,s') = 0.8 intended direction, 0.1 each perpendicular.`
  - If the move hits a wall or Blocked cell, agent stays at s.
- `Rewards: R(*,*,(3,4)) = +1 [Food],  R(*,*,(2,4)) = -1 [Tiger],  R(*,*,*) = -0.2 [all others]`

- Discount factor g = 0.95
- Terminal states (3, 4) = Food and (2, 4) = Tiger are absorbing with V= 0

V^t values used (from Qn)
- Row 3: V(3,1) = 0.197, V(3,2) = 0.570, V(3, 3) = 0.882, V(3, 4) = **Food**=0
- Row 2: V(2,1) = -0.177, V(2,2) = **Blocked**, V(2, 3) = 0.425, V(2, 4) = **Tiger**=0
- Row 1: V(1,1) = -0.548, V(1,2) = -0.291, V(1, 3) = 0.042, V(1, 4)= -0.340

Q(s,a) calc for the three missing states: <FOR CODE VERIFICATION SEE END! >

State (3,3)

[East=Food, North=wall, South=(2,3), West=(3,2)] → Action EAST
0.8 * [+1.000 + 0.95 * 0.000] = +0.8000 // intended: hits Food
0.1 * [-0.200 + 0.95 * 0.882] = +0.0638 // perp N: wall, stays (3,3)
0.1 * [-0.200 + 0.95 * 0.425] = +0.0204 // perp S: moves to (2,3)

Q(3,3,E) = 0.884

Other actions: N=0.644 W=0.357 S=0.297
V^(t+1)(3,3) = 0.884, Policy states = East

State (2,1)

[North=(3,1), South=(1,1), East=Blocked->self, West=wall->self] → Action NORTH
0.8 * [-0.200 + 0.95 * 0.197] = -0.0103 // intended: (3,1)
0.1 * [-0.200 + 0.95 * -0.177] = -0.0368 // perp E: Blocked, stays (2,1)
0.1 * [-0.200 + 0.95 * -0.177] = -0.0368 // perp W: wall, stays (2,1)

Q(2,1,N) = -0.084

Other actions: S=-0.650 E/W=-0.368
V^(t+1)(2,1) = -0.084 policy = North

State (2,3)

[North=(3,3), East=Tiger, West=Blocked->self, South=(1,3)] → Action NORTH
0.8 * [-0.200 + 0.95 * 0.882] = +0.5103 // intended: (3,3)
0.1 * [-1.000 + 0.95 * 0.000] = -0.1000 // perp E: Tiger
0.1 * [-0.200 + 0.95 * 0.425] = +0.0204 // perp W: Blocked, stays (2,3)

Q(2,3,N) = 0.431

Other actions: W=0.211 S=-0.208 E=-0.752
V^(t+1)(2,3) = 0.431 policy = North

Results

V^(t+1):

	C1	C2	C3	C4
R1	0.235	0.579	0.884	Food (0)
R2	-0.084	Blocked	0.431	Tiger (0)
R3	-0.414	-0.223	0.063	-0.280

<FOR CODE VERIFICATION SEE END! >

Assumptions/ Justifications

- Value iteration is a DP for MDPs → Each Bellman backup computes the optimal value by maximising over all actions the expected return + convergence guaranteed for finite MDPs with g<1 as per Sutton and Barto
- when the intended or perpendicular move is blocked, the agent remains at s + Reward for self-transition is R(s,a,s) = -0.2 and value used is V^t(s)
- Terminal states (Food, Tiger) have V^t = V^(t+1) = 0 by definition -- episode ends there so no future rewards accumulate
- g = 0.95 means future rewards are discounted but still significantly weighted so the agent cares about reaching Food even when quite a few steps away

(2) Determine the policy for each state marked with "?" using the Q-function Q^(t+1) (s,a). Fill in the missing actions in the table below and justify your answers based on Q^(t+1) (s,a). [3 points]

Steps

π(s) = argmax_a Q^(t+1)(s, a)
The three '?' states are (3,3), (2,1), and (2,3)

FOR State (3, 3):

1. Q(E)=0.884 Q(N)=0.644 Q(W)=0.357 Q(S)=0.297
2. argmax = East → pi(3,3) = East
3. (3, 3) is immediately adjacent to Food (3,4) **AND** East gives 0.8 prob of +1 reward in one step

FOR State (2, 1):

4. Q(N)=-0.084 Q(E)=Q(W)=-0.368 Q(S)=-0.650
5. argmax = North → pi(2,1) = North
6. (2, 1) is a stuck state with (East=Blocked, West=wall) **AND** North toward (3,1) is the only escape toward Food

FOR State (2, 3):

7. Q(N)=0.431 Q(W)=0.211 Q(S)=-0.208 Q(E)=-0.752
8. argmax = North → pi(2,3) = North
9. North reaches (3,3) which is one step from Food **AND** East risks Tiger as a perpendicular outcome

<FOR CODE VERIFICATION SEE END! >

Results

(π^(t+1)):

	C1	C2	C3	C4
R1	E	E	E	Food (Terminal)
R2	N	Blocked	N	Tiger (Terminal)
R3	N	E	N	W

(3, 3) = East, (2, 1) = North, (2, 3) = North

<FOR CODE VERIFICATION SEE END! >

Assumptions/ Justifications

- Policy is derived from Q^(t+1), not Q^t: question explicitly asks for π using Q^(t+1)(s,a) so the Q-function at iteration t+1, which uses V^t(s') in its computation
- All the ties in Q-values (like if Q(E)=Q(W) for state (2,1)) are handled convergence

CODE VERIFICATION

7 steps

```
GAMMA  = 0.95
STEP_R = -0.2

FOOD    = (2, 3)   # row3, col4
TIGER   = (1, 3)   # row2, col4
BLOCKED = (1, 1)   # row2, col2
TERMINAL = {FOOD, TIGER}
```

Discount factor g = 0.95

Step reward R = -0.2 (all non-terminal transitions)

Food (2, 3) R=+1 terminal, V=0

Tiger (1, 3) R=-1 terminal, V=0

Blocked (1, 1) agent cannot enter

```
####### TABLE
Vt = { (2,0): 0.197, (2,1): 0.570, (2,2): 0.882, (2,3): 0.000, (1,0): -0.177, (1,2): 0.425,
(1,3): 0.000, (0,0): -0.548, (0,1): -0.291, (0,2): 0.042, (0,3): -0.340, }
print("=" * 60)
print("STEP 2 — V^t table (from problem statement)")
print("=" * 60)
print("       col1    col2    col3    col4")
for r in [2, 1, 0]:
    row_label = f"row{r+1}"
    vals = []
    for c in range(4):
        if (r,c) == BLOCKED:
            vals.append("  Block")
        elif (r,c) == FOOD:
            vals.append("Food(0)")
        elif (r,c) == TIGER:
            vals.append(" Tig(0)")
        else:
            vals.append(f"{Vt[(r,c)]:+7.3f}")
    print(f"  {row_label}:  {'  '.join(vals)}")
print()
```

col1 col2 col3 col4

row3: +0.197 +0.570 +0.882 Food(0)

row2: -0.177 Block +0.425 Tig(0)

row1: -0.548 -0.291 +0.042 -0.340

```
####### TRANSITION MODEL
ACTIONS = {'N': (1,0), 'S': (1,0), 'E': (0,1), 'W': (0,-1)}
PERP    = {'N':['E','W'], 'S':['E','W'], 'E':['N','S'], 'W':['N','S']}

def next_state(r, c, action):
    dr, dc = ACTIONS[action]
    nr, nc = r+dr, c+dc
    if nr < 0 or nr >= 3 or nc < 0 or nc >= 4 or (nr,nc) == BLOCKED:
        return r, c   # bounce back
    return nr, nc

def get_reward(r, c, nr, nc):
    if (nr,nc) == FOOD:  return +1.0
            if (nr,nc) == TIGER: return -1.0
            return STEP_R

print("=" * 60)
print("STEP 3 — Transition model verification")
print("=" * 60)
print("  T(s,a,s') = 0.8 intended,  0.1 each perpendicular")
print("  Wall/Blocked -> agent bounces back to s")
print()
```

```
# spot check transitions for the 3 missing states
for label,(r,c),action in [("(3,3)",(2,2),'E'),
                           ("(2,1)",(1,0),'N'),
                           ("(2,3)",(1,2),'N')]:
    p1 = PERP[action]
    s_i  = next_state(r,c,action)
    s_p1 = next_state(r,c,p1[0])
    s_p2 = next_state(r,c,p1[1])
    print(f"  State {label} action {action}:")
    print(f"    0.8  intended {action}  -> {s_i}")
    print(f"    0.1  perp    {p1[0]}  -> {s_p1}")
    print(f"    0.1  perp    {p1[1]}  -> {s_p2}")
    print()
```

T(s,a,s') = 0.8 intended, 0.1 each perpendicular

Wall/Blocked -> agent bounces back to s

State (3,3) action E:

0.8 intended E -> (2, 3)

0.1 perp   N -> (2, 2)

0.1 perp   S -> (1, 2)

State (2,1) action N:

0.8 intended N -> (2, 0)

0.1 perp   E -> (1, 0)

0.1 perp   W -> (1, 0)

State (2,3) action N:

0.8 intended N -> (2, 2)

0.1 perp   E -> (1, 3)

0.1 perp   W -> (1, 2)

```
# Q FUNCTION
def Q(r, c, action, verbose=False):
    outcomes = [
        (0.8, action,"intended"),
        (0.1, PERP[action][0], f"perp {PERP[action][0]}"),
        (0.1, PERP[action][1], f"perp {PERP[action][1]}"),
    ]
    total = 0.0
    if verbose:
        print(f"    {'T':>4}  {'s_next':>8}  {'R':>7}  {'V^t':>7}  {'T*[R+g*V]':>10}")
        print(f"    {'—'*50}")
    for T, a, desc in outcomes:
        nr, nc  = next_state(r, c, a)
        R       = get_reward(r, c, nr, nc)
        Vt_next = Vt.get((nr,nc), 0.0)
        contrib = T * (R + GAMMA * Vt_next)
        total  += contrib
        if verbose:
            print(f"    {T:>4.1f}  ({nr+1},{nc+1}){' ':>4}  "
                  f"{R:>+7.3f}  {Vt_next:>+7.3f}  {contrib:>+10.4f}"
                  f"   // {desc}")
    if verbose:
        print(f"    {'—'*50}")
        print(f"    {'Q =':>38}  {total:>+10.4f}")
    return total
```

```
============================================================
STEP 4 — Detailed Q(s,a) proof for 3 missing states
============================================================
  ——— State (3,3) ——————————————————————————————————————————
  Action E (best):
```

```
   T    s_next       R      V^t   T*[R+g*V]
 ──────────────────────────────────────────────────────────────
  0.8  (3,4)      +1.000   +0.000    +0.8000   // intended
  0.1  (3,3)      -0.200   +0.882    +0.0638   // perp N
  0.1  (2,3)      -0.200   +0.425    +0.0204   // perp S
 ──────────────────────────────────────────────────────────────
                          Q =     +0.8842
All Q values:
  Q((3,3),N) = +0.6445
  Q((3,3),S) = +0.2972
  Q((3,3),E) = +0.8842 <-- max
  Q((3,3),W) = +0.3574
V^(t+1)(3,3) = +0.8842  (expected +0.884)  PASS ✓
policy (3,3)  = E        (expected E)      PASS ✓

──── State (2,1) ──────────────────────────────────────────────────────
Action N (best):
   T    s_next       R      V^t   T*[R+g*V]
 ──────────────────────────────────────────────────────────────
  0.8  (3,1)      -0.200   +0.197    -0.0103   // intended
  0.1  (2,1)      -0.200   -0.177    -0.0368   // perp E
  0.1  (2,1)      -0.200   -0.177    -0.0368   // perp W
 ──────────────────────────────────────────────────────────────
                          Q =     -0.0839
All Q values:
  Q((2,1),N) = -0.0839 <-- max
  Q((2,1),S) = -0.6501
  Q((2,1),E) = -0.3679
  Q((2,1),W) = -0.3679

V^(t+1)(2,1) = -0.0839  (expected -0.084)  PASS ✓
policy (2,1)  = N        (expected N)      PASS ✓

──── State (2,3) ──────────────────────────────────────────────────────
Action N (best):
   T    s_next       R      V^t   T*[R+g*V]
 ──────────────────────────────────────────────────────────────
  0.8  (3,3)      -0.200   +0.882    +0.5103   // intended
  0.1  (2,4)      -1.000   +0.000    -0.1000   // perp E
  0.1  (2,3)      -0.200   +0.425    +0.0204   // perp W
 ──────────────────────────────────────────────────────────────
                          Q =     +0.4307
All Q values:
  Q((2,3),N) = +0.4307 <-- max
  Q((2,3),S) = -0.2077
  Q((2,3),E) = -0.7522
  Q((2,3),W) = +0.2108
V^(t+1)(2,3) = +0.4307  (expected +0.431)  PASS ✓
policy (2,3)  = N        (expected N)      PASS ✓
```

```
##### THE. FULL GRID V^(t+1)
print()
print("=" * 60)
print("STEP 5 — Full V^(t+1) grid")
print("=" * 60)

expected_V = {
   (2,0): 0.235,  (2,1): 0.579,  (2,2): 0.884,  (2,3): 0.000,
   (1,0):-0.084,                 (1,2): 0.431,  (1,3): 0.000,
   (0,0):-0.414,  (0,1):-0.223,  (0,2): 0.063,  (0,3):-0.280,
}

Vt1 = {}
for r in range(3):
   for c in range(4):
      if (r,c) == BLOCKED:      Vt1[(r,c)] = None
      elif (r,c) in TERMINAL:   Vt1[(r,c)] = 0.0
      else:
         qs = {a: Q(r,c,a) for a in ACTIONS}
         Vt1[(r,c)] = max(qs.values())

print()
print("computed V^(t+1):")
print("       col1    col2    col3    col4")
for r in [2,1,0]:
   vals = []
   for c in range(4):
      if (r,c) == BLOCKED:    vals.append("  Block")
      elif Vt1[(r,c)] == 0.0: vals.append(f"  0.000")
      else:                   vals.append(f"{Vt1[(r,c)]:+7.3f}")
   print(f"  row{r+1}:  {'  '.join(vals)}")

print()
print("  Verification against expected values:")
print(f"  {'State':<8} {'Computed':>10} {'Expected':>10} {'Diff':>8} {'Pass?':>8}")
print(f"  {'—'*48}")
all_pass = True
for (r,c), exp in sorted(expected_V.items(), reverse=True):
   if (r,c) == BLOCKED: continue
   got  = Vt1[(r,c)]
   diff = abs(got - exp)
   ok   = diff < 0.002
   if not ok: all_pass = False
   print(f"  ({r+1},{c+1})     {got:>+10.4f} {exp:>+10.3f} {diff:>8.5f} "
         f"{'PASS ✓' if ok else 'FAIL ×'}")

print()
print(f"  All values correct: {'YES ✓' if all_pass else 'NO ×'}")
```

```
============================================================
STEP 5 — Full V^(t+1) grid
============================================================

computed V^(t+1):
       col1    col2    col3    col4
  row3:  +0.235  +0.579  +0.884   0.000
  row2:  -0.084  Block  +0.431   0.000
  row1:  -0.414  -0.223  +0.063  -0.280

  Verification against expected values:
  State     Computed   Expected     Diff    Pass?
  ————————————————————————————————————————————————————————————————————————————————

  (3,4)       +0.0000     +0.000  0.00000 PASS ✓
```

```
(3,3)      +0.8842    +0.884  0.00017 PASS ✓
(3,2)      +0.5786    +0.579  0.00038 PASS ✓
(3,1)      +0.2351    +0.235  0.00010 PASS ✓
(2,4)      +0.0000    +0.000  0.00000 PASS ✓
(2,3)      +0.4307    +0.431  0.00030 PASS ✓
(2,1)      -0.0839    -0.084  0.00009 PASS ✓
(1,4)      -0.2804    -0.280  0.00038 PASS ✓
(1,3)      +0.0631    +0.063  0.00005 PASS ✓
(1,2)      -0.2234    -0.223  0.00037 PASS ✓
(1,1)      -0.4142    -0.414  0.00023 PASS ✓

All values correct: YES ✓
```

```python
#### FULL POLICY TABLE
print()
print("=" * 60)
print("STEP 6 — Full policy table pi^(t+1)")
print("=" * 60)

expected_pi = {
    (2,0):'E', (2,1):'E', (2,2):'E',
    (1,0):'N',            (1,2):'N',
    (0,0):'N', (0,1):'E', (0,2):'N', (0,3):'W',
}

Pi = {}
for r in range(3):
    for c in range(4):
        if (r,c) in TERMINAL or (r,c) == BLOCKED:
            Pi[(r,c)] = '-'
        else:
            qs = {a: Q(r,c,a) for a in ACTIONS}
            Pi[(r,c)] = max(qs, key=qs.get)

print()
print("  Computed policy:")
print("       col1   col2   col3   col4")
for r in [2,1,0]:
    vals = []
    for c in range(4):
        if   (r,c) == BLOCKED:vals.append("  Blk")
        elif (r,c) == FOOD: vals.append(" Food")
        elif (r,c) == TIGER: vals.append(" Tig ")
        else: vals.append(f"    {Pi[(r,c)]}")
    print(f"  row{r+1}:  {'  '.join(vals)}")

print()
print("  Verification against expected policy:")
print(f"  {'State':<8} {'Computed':>10} {'Expected':>10} {'Pass?':>8}")
print(f"  {'—'*40}")
all_pi_pass = True
for (r,c), exp in sorted(expected_pi.items(), reverse=True):
    got = Pi[(r,c)]
    ok  = got == exp
    if not ok: all_pi_pass = False
    print(f"  ({r+1},{c+1})     {got:>10} {exp:>10} {'PASS ✓' if ok else 'FAIL ✗'}")

print()
        print(f"  All policies correct: {'YES ✓' if all_pi_pass else 'NO ✗'}")
        print()
        print("=" * 60)
        print(f"FINAL: All V^(t+1) correct: {'YES ✓' if all_pass else 'NO ✗'}")
        print(f"FINAL: All pi correct:      {'YES ✓' if all_pi_pass else 'NO ✗'}")
        print("=" * 60)
```

```
============================================================
STEP 6 — Full policy table pi^(t+1)
============================================================

 Computed policy:
      col1   col2   col3   col4
 row3:     E      E      E   Food
 row2:     N    Blk      N   Tig
 row1:     N      E      N      W

 Verification against expected policy:
 State      Computed   Expected    Pass?
 ──────────────────────────────────────────────────────────────────────

 (3,3)            E         E PASS ✓
 (3,2)            E         E PASS ✓
 (3,1)            E         E PASS ✓
 (2,3)            N         N PASS ✓
 (2,1)            N         N PASS ✓
 (1,4)            W         W PASS ✓
 (1,3)            N         N PASS ✓
 (1,2)            E         E PASS ✓
 (1,1)            N         N PASS ✓

 All policies correct: YES ✓


============================================================
FINAL: All V^(t+1) correct: YES ✓
FINAL: All pi correct:       YES ✓
============================================================
```